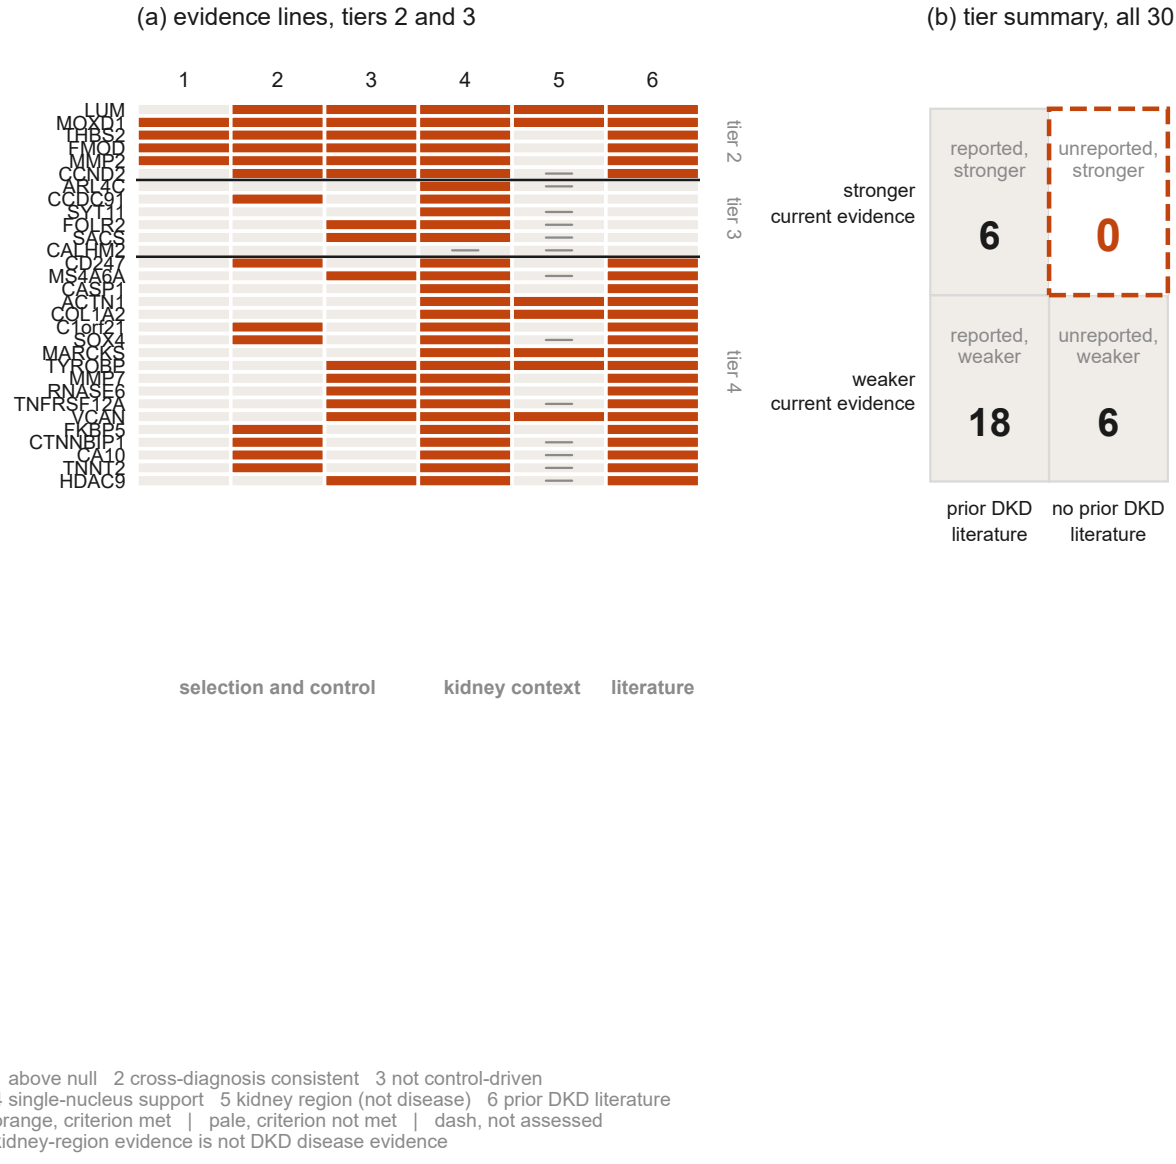
Additional file 5. Candidate evidence, all 30 candidates

The main text describes the 12 candidates in tiers 2 and 3 – those that are either strong in the data or absent from the DKD literature. This file shows the evidence matrix over all 30 specificity-gated candidates, adding the 18 in tier 4, which are neither. Nothing is computed differently here; only the row selection differs.



1 above null 2 cross-diagnosis consistent 3 not control-driven
4 single-nucleus support 5 kidney region (not disease) 6 prior DKD literature
orange, criterion met | pale, criterion not met | dash, not assessed
kidney-region evidence is not DKD disease evidence

Panel **a**: a filled cell means the evidence line is met, an open cell that it was tested and not met, and a hatched cell that the gene could not be assessed on that line. Single-nucleus and proteomic corroboration are not available for every gene, and scoring an absent measurement as a failure would confuse the two. Panel **b** counts all 30 candidates in both versions of the figure, so the two panels always describe the same set. Tier assignments are read from the candidate table rather than recomputed here.